

# 高性能report2

## 高性能计算编程实验2

### 1.实验环境

CPU：Dell Inc. OptiPlex Micro Plus 7020 Intel® Core™ i9-14900 × 32 64g

Os：ubuntu 22.04.05 LTS

Compiler: g++ (Ubuntu 11.4.0)

### 2.修改代码为SoA分配方式

主要修改原代码使用的数据结构XYZ和Element. 原代码中声明单一变量组成的结构体，然后通过结构体数组的形式存储数据。修改后，在声明时就确定结构体的成员为数组。

```
typedef struct
{
    double* x;
    double* y;
    double* z;
} XYZ_SoA;

typedef struct
{
    double* mass;
    XYZ_SoA acct;
    XYZ_SoA noused; // but here
    XYZ_SoA Velocity;
} Element;
```

编译：

```
g++ -g -O2 ./homework2.cpp -o hw2
```

运行：

```
numactl --membind=0 --cpubind=0 ./hw2
```

运行情况：

```
hnm@hnm-OptiPlex-Micro-Plus-7020:~/code/hiper2/hiper2$ g++ -g -O2 ./homework2.cpp -o hw2
hnm@hnm-OptiPlex-Micro-Plus-7020:~/code/hiper2/hiper2$ numactl --membind=0 --cpubind=0 ./hw2
muladd,timing=685779us
sum = 5.75430668e+09,timing=55259us
```

优化后运行情况：

```
root@hbm-OptiPlex-Micro-Plus-7020:/home/hhm/code/hiper2/hiper2# numactl --membind=0 --cpubind=0 ./hw2_SoA
muladd,timing=248724us
sum = 5.75430668e+09,timing=18205us
```

### 3. SoA 优化

#### a. 内存对齐

使用 `std::aligned_alloc` 来确保每个数组（如 `elements.mass`、`elements.acct.x` 等）是按 64 字节对齐的。

```
elements.mass = (double*)std::aligned_alloc(64, ELEMENT_NUM * sizeof(double)); // 64-byte aligned
elements.acct.x = (double*)std::aligned_alloc(64, ELEMENT_NUM * sizeof(double));
elements.acct.y = (double*)std::aligned_alloc(64, ELEMENT_NUM * sizeof(double));
elements.acct.z = (double*)std::aligned_alloc(64, ELEMENT_NUM * sizeof(double));
elements.Velocity.x = (double*)std::aligned_alloc(64, ELEMENT_NUM * sizeof(double));
elements.Velocity.y = (double*)std::aligned_alloc(64, ELEMENT_NUM * sizeof(double));
elements.Velocity.z = (double*)std::aligned_alloc(64, ELEMENT_NUM * sizeof(double));
elements.noused.x = (double*)std::aligned_alloc(64, ELEMENT_NUM * sizeof(double));
elements.noused.y = (double*)std::aligned_alloc(64, ELEMENT_NUM * sizeof(double));
elements.noused.z = (double*)std::aligned_alloc(64, ELEMENT_NUM * sizeof(double));
```

优化后运行结果：

```
root@hbm-OptiPlex-Micro-Plus-7020:/home/hhm/code/hiper2/hiper2# numactl --membind=0 --cpubind=0 ./al
muladd,timing=248112us
sum = 5.75430668e+09,timing=18272us
```

#### b. 循环展开

通过将每次更新拆分为 4 次操作，尝试减少循环控制的开销

```
for (; ii + 3 < ELEMENT_NUM; ii += 4)
{
    elements.Velocity.x[ii] += dt * elements.acct.x[ii];
    elements.Velocity.x[ii + 1] += dt * elements.acct.x[ii + 1];
    elements.Velocity.x[ii + 2] += dt * elements.acct.x[ii + 2];
    elements.Velocity.x[ii + 3] += dt * elements.acct.x[ii + 3];

    elements.Velocity.y[ii] += dt * elements.acct.y[ii];
    elements.Velocity.y[ii + 1] += dt * elements.acct.y[ii + 1];
    elements.Velocity.y[ii + 2] += dt * elements.acct.y[ii + 2];
    elements.Velocity.y[ii + 3] += dt * elements.acct.y[ii + 3];

    elements.Velocity.z[ii] += dt * elements.acct.z[ii];
    elements.Velocity.z[ii + 1] += dt * elements.acct.z[ii + 1];
    elements.Velocity.z[ii + 2] += dt * elements.acct.z[ii + 2];
    elements.Velocity.z[ii + 3] += dt * elements.acct.z[ii + 3];
}

// 处理剩余部分
```

运行结果：

```
root@hbm-OptiPlex-Micro-Plus-7020:/home/hbm/code/hiper2/hiper2# numactl --membind=0 --cpubind=0 ./lr
muladd,timing=246834us
sum = 5.75430668e+09,timing=18012us
```

完全不使用循环展开：

```
root@hbm-OptiPlex-Micro-Plus-7020:/home/hbm/code/hiper2/hiper2# numactl --membind=0 --cpubind=0 ./no_lr_test
muladd,timing=255908us
sum = 5.75430668e+09,timing=18109us
```

## c.fma

通过fma将乘法和加法合并成一条指令

```
for (size_t step = 0; step < 10; ++step)
{
    for (size_t ii = 0; ii < ELEMENT_NUM; ++ii)
    {
        elements.Velocity.x[ii] = fma(dt, elements.acct.x[ii], elements.Velocity.x[ii]);
        elements.Velocity.y[ii] = fma(dt, elements.acct.y[ii], elements.Velocity.y[ii]);
        elements.Velocity.z[ii] = fma(dt, elements.acct.z[ii], elements.Velocity.z[ii]);
    }
    dt = dt * (((rand() % 10) / 10.0) * 2.0);
}
size_t finish = GetUsec();
printf("muladd,timing=%ldus\n", finish - start);

start = GetUsec();

// 使用 FMA 优化代码
double sum = 0.0;
for (size_t ii = 0; ii < ELEMENT_NUM; ++ii)
{
    // fma(a, b, c) 表示 a * b + c
    sum += 0.5 * elements.mass[ii] *
        (fma(elements.Velocity.x[ii], elements.Velocity.x[ii], fma(elements.Velocity.y[ii], elements.Velocity.y[ii],
            fma(elements.Velocity.z[ii], elements.Velocity.z[ii], 0.0)))));
}
```

运行结果：

```
root@hbm-OptiPlex-Micro-Plus-7020:/home/hbm/code/hiper2/hiper2# numactl --membind=0 --cpubind=0 ./fma
muladd,timing=238523us
sum = 5.75430668e+09,timing=15693us
```

## 4.结论

从运行时间来看，AoS改为SoA对代码的运行时间有较大改善，这是由于SoA分配方式减少了大量的寻址工作导致的。使用perf stat对两个程序进行分析后发现SoA程序的IPC高于原程序，说明其指令流水线更加高效。此外，SoA的LLC访问减少，说明减少了高延迟的主存访问。

```

3,921,978,494      cpu_atom/instructions/      #    0.45  insn per cycle      (0.06%)
13,494,317,404    cpu_core/instructions/    #    1.53  insn per cycle      (99.94%)
263,474,962       cpu_atom/branches/        # 124.686 M/sec      (0.06%)
2,427,675,533     cpu_core/branches/         #    1.149 G/sec      (99.94%)
   647,472        cpu_atom/branch-misses/    #    0.25% of all branches  (0.06%)
  4,103,339        cpu_core/branch-misses/    #    1.56% of all branches  (99.94%)
  TopdownL1 (cpu_core)      #  70.3 % tma_backend_bound
                          #    0.9 % tma_bad_speculation
                          #    4.8 % tma_frontend_bound
                          #  24.0 % tma_retiring      (99.94%)
  TopdownL1 (cpu_atom)      #   -0.3 % tma_bad_speculation
                          #    8.7 % tma_retiring      (0.06%)
                          #   91.1 % tma_backend_bound
                          #   91.1 % tma_backend_bound_aux
                          #    0.5 % tma_frontend_bound  (0.06%)

<not counted>      L1-dcache-loads      (0.00%)
3,617,437,695      L1-dcache-loads      #    1.712 G/sec      (99.94%)
<not supported>    L1-dcache-load-misses
  56,870,151        L1-dcache-load-misses  (99.94%)
<not counted>      LLC-loads      (0.00%)
 114,320,175        LLC-loads      #   54.100 M/sec      (99.94%)
<not counted>      LLC-load-misses  (0.00%)
 113,477,263        LLC-load-misses  (99.94%)

2.122475415 seconds time elapsed

1.851734000 seconds user
0.262678000 seconds sys

root@hbm-OptiPlex-Micro-Plus-7020:/home/hbm/code/hiper2/hiper2# perf stat -d ./hw2_SoA
muladd,timing=248946us
sum = 5.75430668e+09,timing=18399us

Performance counter stats for './hw2_SoA':

  1,574.71 msec task-clock      #    0.999 CPUs utilized
      23      context-switches  #   14.606 /sec
      6      cpu-migrations     #    3.810 /sec
  229,503     page-faults      # 145.743 K/sec
  6,583,351,352  cpu_atom/cycles/      #    4.181 GHz      (0.03%)
  8,919,562,038  cpu_core/cycles/      #    5.664 GHz      (99.97%)
  7,015,955,730  cpu_atom/instructions/ #    1.07  insn per cycle  (0.03%)
12,497,904,768  cpu_core/instructions/ #    1.90  insn per cycle  (99.97%)
  1,467,081,696  cpu_atom/branches/    #  931.652 M/sec      (0.03%)
  2,318,178,342  cpu_core/branches/    #    1.472 G/sec      (99.97%)
   4,923,778     cpu_atom/branch-misses/ #    0.34% of all branches  (0.03%)
   4,245,399     cpu_core/branch-misses/ #    0.29% of all branches  (99.97%)
  TopdownL1 (cpu_core)      #  65.5 % tma_backend_bound
                          #    0.8 % tma_bad_speculation
                          #    5.5 % tma_frontend_bound
                          #  28.2 % tma_retiring      (99.97%)
  TopdownL1 (cpu_atom)      #    7.4 % tma_bad_speculation
                          #   26.2 % tma_retiring      (0.03%)
                          #   61.1 % tma_backend_bound
                          #   61.1 % tma_backend_bound_aux
                          #    5.3 % tma_frontend_bound  (0.03%)

<not counted>      L1-dcache-loads      (0.00%)
 3,534,487,584      L1-dcache-loads      #    2.245 G/sec      (99.97%)
<not supported>    L1-dcache-load-misses
  90,452,602        L1-dcache-load-misses  (99.97%)
<not counted>      LLC-loads      (0.00%)
   7,216,977        LLC-loads      #    4.583 M/sec      (99.97%)
<not counted>      LLC-load-misses  (0.00%)
   6,405,468        LLC-load-misses  (99.97%)

1.576708947 seconds time elapsed

1.375940000 seconds user
0.200136000 seconds sys

```

然而，后面采用的几种优化方法似乎都没有达到预期的效果。通过添加选项禁止向量化和循环展开，确认并不是-O2编译级别中包含向量化和循环展开，而是以上三种发放对于原代码都不是特别有效。其中一种可能是，手动的循环展开只有4次，因此观察不到太大的效果。另外，SoA分配方式所申请的内存是连续的内存块，已经保证了内存空间的连续分布，因此可能不需要内存对齐。

```

root@hbm-OptiPlex-Micro-Plus-7020:/home/hbm/code/hiper2/hiper2# perf stat -d ./lr
muladd,timing=250923us
sum = 5.75430668e+09,timing=17491us

```

Performance counter stats for './lr':

1,622.74 msec	task-clock	#	0.996 CPUs utilized	
96	context-switches	#	59.159 /sec	
9	cpu-migrations	#	5.546 /sec	
229,505	page-faults	#	141.431 K/sec	
<not counted>	cpu_atom/cycles/			(0.00%)
8,911,902,844	cpu_core/cycles/	#	5.492 GHz	
<not counted>	cpu_atom/instructions/			(0.00%)
12,213,633,959	cpu_core/instructions/			
<not counted>	cpu_atom/branches/			(0.00%)
2,190,284,985	cpu_core/branches/	#	1.350 G/sec	
<not counted>	cpu_atom/branch-misses/			(0.00%)
4,242,179	cpu_core/branch-misses/			
TopdownL1 (cpu_core)		#	68.2 % tma_backend_bound	
		#	1.2 % tma_bad_speculation	
		#	4.0 % tma_frontend_bound	
		#	26.6 % tma_retiring	
<not counted>	L1-dcache-loads			(0.00%)
3,531,179,371	L1-dcache-loads	#	2.176 G/sec	
<not supported>	L1-dcache-load-misses			
68,304,949	L1-dcache-load-misses			
<not counted>	LLC-loads			(0.00%)
7,910,587	LLC-loads	#	4.875 M/sec	
<not counted>	LLC-load-misses			(0.00%)
6,539,999	LLC-load-misses			
1.629042372 seconds time elapsed				
1.414791000 seconds user				
0.209524000 seconds sys				

```

root@hbm-OptiPlex-Micro-Plus-7020:/home/hbm/code/hiper2/hiper2# perf stat -d ./fma
muladd,timing=255223us
sum = 5.75430668e+09,timing=18068us

```

Performance counter stats for './fma':

1,642.36 msec	task-clock	#	0.999 CPUs utilized	
67	context-switches	#	40.795 /sec	
3	cpu-migrations	#	1.827 /sec	
229,504	page-faults	#	139.740 K/sec	
<not counted>	cpu_atom/cycles/			(0.00%)
8,875,573,789	cpu_core/cycles/	#	5.404 GHz	
<not counted>	cpu_atom/instructions/			(0.00%)
12,482,991,768	cpu_core/instructions/			
<not counted>	cpu_atom/branches/			(0.00%)
2,316,355,804	cpu_core/branches/	#	1.410 G/sec	
<not counted>	cpu_atom/branch-misses/			(0.00%)
4,239,335	cpu_core/branch-misses/			
TopdownL1 (cpu_core)		#	66.1 % tma_backend_bound	
		#	1.2 % tma_bad_speculation	
		#	4.7 % tma_frontend_bound	
		#	28.0 % tma_retiring	
<not counted>	L1-dcache-loads			(0.00%)
3,531,480,727	L1-dcache-loads	#	2.150 G/sec	
<not supported>	L1-dcache-load-misses			
87,969,207	L1-dcache-load-misses			
<not counted>	LLC-loads			(0.00%)
7,005,302	LLC-loads	#	4.265 M/sec	
<not counted>	LLC-load-misses			(0.00%)
5,967,763	LLC-load-misses			
1.643961483 seconds time elapsed				